

Adem Mrabet

Tunis, Tunisia (open to relocating to Berlin) | adem.mrabet99@gmail.com | +216 58 407 600 | linkedin.com/in/ademmrabet | github.com/ademmrabet

September 22, 2026

Dear Almedia Hiring Team,

I'm writing to apply for the AI Engineer – Data role. What drew me in wasn't the job title, it was the brief: replace a rigid reward design system with something genuinely agentic — AI that looks into the games, generates offers, evaluates them, and only then goes live. That evaluate-before-ship loop is exactly the pattern I've already built and shipped, and I'd like to bring it to your rewards engine.

For my final-year AI engineering project at the African Development Bank, I was handed a problem with the same shape as the one you describe: a fuzzy “could AI help here?” question, no existing dataset, and a business-critical process (who is authorized to approve what, and under which conditions) that couldn't tolerate a wrong answer. I built the whole system myself — an 8-stage pipeline that turned an unstructured 250-page governance document into a 327-node knowledge graph, a hybrid retrieval layer, and a grounded-generation step where an LLM only rephrases pre-assembled facts. The part I'm proudest of is the verifier sitting on top: every rephrased answer is checked against the source data, and anything that drifts gets discarded before a user ever sees it. It's a small-scale version of the eval-driven, agentic pattern you're describing for reward design, and it taught me firsthand where to trust an LLM and where not to.

That project shipped as a full production platform (FastAPI, React, PostgreSQL, Docker) that I deployed and still operate, and I've since built a second RAG system (FAISS + BERT embeddings + Groq generation) and a distributed pipeline over large-scale transactional data with PySpark. I code in Python daily, I'm comfortable moving from a vague brief to a working system without waiting to be told how, and I'd rather ship a small test against real data than polish a plan. I also know when not to reach for AI: several of the simpler, rule-based pieces of my pipeline exist because a deterministic check was more reliable than a model call, and I chose that on purpose.

I'm based in Tunisia and ready to relocate to Berlin — I understand this is an in-person, five-days-a-week role, and I'd welcome that. I'd love the chance to talk about how I could help build the first dedicated AI-engineering patterns for your data team.

Thank you for your time and consideration.

Best regards,

Adem Mrabet